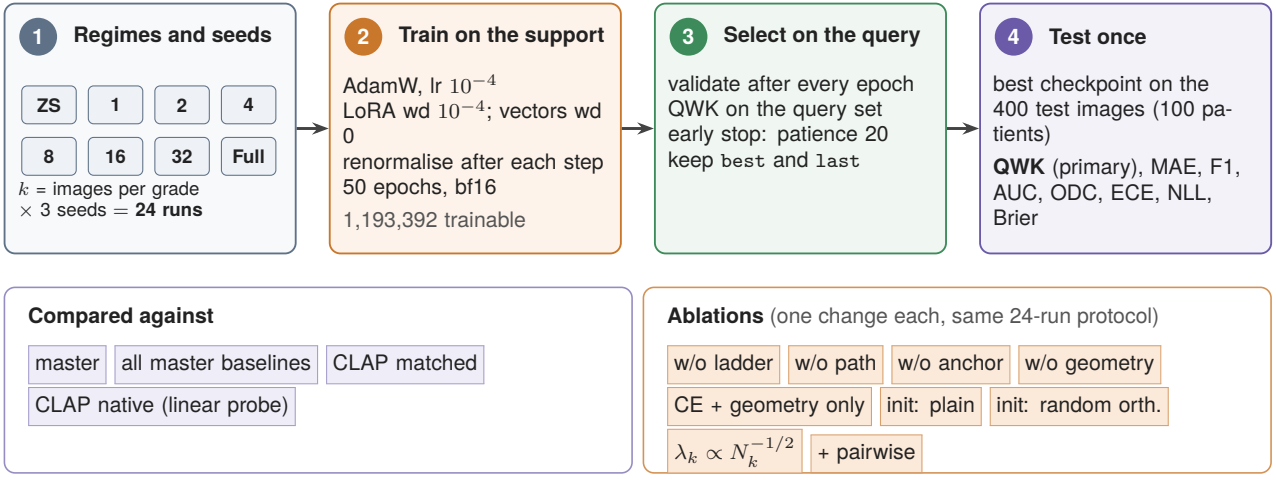
1
Regimes and seeds
ZS
1
2
4
8
16
32
Full
k = images per grade
× 3 seeds = 24 runs
2
Train on the support
AdamW, lr 10^{-4}
LoRA wd 10^{-4}; vectors wd 0
renormalise after each step
50 epochs, bf16
1,193,392 trainable
3
Select on the query
validate after every epoch
QWK on the query set
early stop: patience 20
keep best and last
4
Test once
best checkpoint on the 400 test images (100 patients)
QWK (primary), MAE, F1, AUC, ODC, ECE, NLL, Brier
Compared against
master
all master baselines
CLAP matched
CLAP native (linear probe)
Ablations (one change each, same 24-run protocol)
w/o ladder
w/o path
w/o anchor
w/o geometry
CE + geometry only
init: plain
init: random orth.
$\lambda_k \propto N_k^{-1/2}$
+ pairwise